

TASK 2 - AWS EBS _ SNAPSHOT_TRANSFER OF VOLUME DATAS

AWS EBS snapshot transfer — step-by-step record with commands, errors & fixes

Goal: Move data from *Instance-1* to *Instance-2* by creating an extra EBS volume on Instance- 1, filling it with files, taking a snapshot, creating a volume from that snapshot and attaching it to Instance- 2 — with full command-by-command notes, explanations, and the errors encountered during your run plus how you resolved them.

Environment / assumptions

- All resources (instances, volumes, snapshots) are in the **same AWS Region**.
 - When creating a volume from a snapshot you must choose the **Availability Zone (AZ)** that contains the instance you will attach the volume to.
 - You have SSH access (pem key) to both instances.
-

Linear step-by-step log (what I did + commands)

This is a linear narrative

Step 1 — Launch EC2 Instance-1 (console)

- Launch a Linux instance (Amazon Linux 2 or Ubuntu). Choose an SSH key pair you control and note the subnet/AZ.

Why: we need a running Linux box to attach and use the extra EBS volume.

Step 2 — Create and attach a 10 GiB EBS volume (Console)

- Console: EC2 → Volumes → Create volume → Size: **10 GiB**, Type: **gp3**, AZ: *match Instance-1 AZ* → Create
- Then: Actions → Attach volume → select **Instance-1** → Device name `/dev/sdf` (console-facing name)

Why: this creates the extra block device that will hold your files.

TASK 2 - AWS EBS _ SNAPSHOT_TRANSFER OF VOLUME DATAS

Step 3 — On Instance-1: identify the new disk

List block devices and filesystems

```
lsblk -f
```

```
dmesg | tail -n 20
```

Map NVMe names to AWS volume IDs if needed

```
ls -l /dev/disk/by-id/
```

Why & expected: `lsblk -f` shows devices, sizes, FSTYPE and mount points. On Nitro instances the AWS-attached device `/dev/sdf` will usually appear in Linux as `/dev/nvme1n1`. Example `lsblk -f` snippet:

NAME	FSTYPE	LABEL	UUID	MOUNTPOINT
nvme0n1				
└─nvme0n1p1	ext4		1111-2222...	/
nvme1n1	ext4		aaaa-bbbb...	

Step 4 — Format the new volume (only on a new/empty volume) and mount it

❑ **Important:** `mkfs` erases any existing data. Only run if the volume is brand new or you intend to wipe it.

Format (one-time)

```
sudo mkfs -t ext4 /dev/nvme1n1
```

create mount point and mount

```
sudo mkdir -p /data
```

```
sudo mount /dev/nvme1n1 /data
```

verify it mounted

```
lsblk -f
```

```
df -h /data
```

Why: `mkfs -t ext4` creates an ext4 filesystem. `mount` attaches it to the filesystem tree so you can read/write under `/data`.

TASK 2 - AWS EBS _ SNAPSHOT_TRANSFER OF VOLUME DATAS

Step 5 — Persist mount across reboots using /etc/fstab (UUID method)

Why: device names (like /dev/nvme1n1) can change; UUID uniquely identifies the filesystem. nofail prevents boot failures when the disk is absent.

```
# find the UUID
sudo blkid -s UUID -o value /dev/nvme1n1
# copy the printed UUID, then add to /etc/fstab (replace <UUID> with the real one)
echo 'UUID=<UUID> /data ext4 defaults,nofail 0 2' | sudo tee -a /etc/fstab

# back up fstab before changes (recommended)
sudo cp /etc/fstab /etc/fstab.bak.$(date +%F-%H%M)

# reload systemd and test mounting
sudo systemctl daemon-reload
sudo umount /data
sudo mount -a

# verify
df -h /data
```

Why: blkid prints the filesystem UUID. echo ... | sudo tee -a /etc/fstab appends the correct fstab line (note the space after -a — **important**). systemctl daemon-reload tells systemd to reread config after modifications.

Step 6 — Create test files on the 10G volume

```
# create test files
sudo mkdir -p /data/docs
echo "Hello from Instance-1" | sudo tee /data/hello.txt
echo "This will move via snapshot" | sudo tee /data/docs/readme.txt
sudo sync
ls -l /data /data/docs
```

Why: tee writes the text into files (with sudo to write to /data), sync flushes caches to disk to make snapshot capture cleaner.

Step 7 — Unmount the volume cleanly (recommended before snapshot)

```
sudo umount /data
# if it's busy, find processes holding it
sudo lsof +f -- /data
# or
sudo fuser -m /data
# then stop the processes and try again
```

TASK 2 - AWS EBS _ SNAPSHOT _ TRANSFER OF VOLUME DATAS

Why: unmounting ensures the filesystem is in a consistent state before creating a snapshot. If you create snapshots of mounted volumes, data may still be OK for most cases, but unmounting is safest.

Step 8 — Create a snapshot of the volume (Console)

- Console: EC2 → Volumes → select your 10 GiB volume → Actions → Create snapshot
- Add a Name/Description: e.g. demo-snap-from-instance1
- Wait until snapshot status = **completed**

Why: a snapshot stores the volume's blocks in S3 and becomes the source to create a new volume later.

Step 9 — Detach original volume from Instance-1 (Console)

- Console: EC2 → Volumes → select the 10 GiB volume → Actions → Detach volume

Why: you want to show transfer by creating a fresh volume for Instance-2 from the snapshot. Detaching prevents accidental concurrent writes.

Step 10 — Launch Instance-2 (if not already running)

- EC2 → Instances → Launch instance → choose Linux (Amazon Linux 2 or Ubuntu), use same VPC and a key pair (so you can SSH in)

Why: Instance-2 is the target where we'll attach the volume created from the snapshot.

Step 11 — Create a new 10 GiB volume from the snapshot in the AZ of Instance-2 (Console)

- EC2 → Snapshots → select snapshot demo-snap-from-instance1 → Actions → Create volume
- Set **Availability Zone** to **match Instance-2**
- Size: 10 GiB; Type: gp3; Create

Why: a volume can only be attached to an instance in the same AZ.

Step 12 — Attach the snapshot-based volume to Instance-2 (Console)

- EC2 → Volumes → select the new volume → Actions → Attach volume → choose Instance-2 → device /dev/sdf
- Wait until status shows **in-use** / attached

Why: now Instance-2 can access the same data that lived on Instance-1.

TASK 2 - AWS EBS _ SNAPSHOT _ TRANSFER OF VOLUME DATAS

Step 13 — On Instance-2: mount and verify files

SSH to Instance-2 and run:

```
lsblk -f
sudo mkdir -p /data
sudo mount /dev/nvme1n1 /data # check 'lsblk' to confirm device name

# Verify mounted size and files
df -h /data
ls -l /data
ls -l /data/docs
cat /data/hello.txt
cat /data/docs/readme.txt
```

Why: you should now see the same files created on Instance-1. This proves the snapshot → new-volume → attach flow worked.

Step 14 — (Optional) Persist mount on Instance-2

If you want the volume to mount automatically on Instance-2 reboot, repeat the UUID→/etc/fstab steps on Instance-2:

```
sudo blkid -s UUID -o value /dev/nvme1n1
# add to /etc/fstab replacing <UUID>
echo 'UUID=<UUID> /data ext4 defaults,nofail 0 2' | sudo tee -a /etc/fstab
sudo systemctl daemon-reload
sudo umount /data
sudo mount -a
```

Step 15 — (Optional) Cleanup when done

- Delete the test volumes and snapshots if you don't want to incur charges: EC2 → Volumes / Snapshots → Delete.

TASK 2 - AWS EBS _ SNAPSHOT_TRANSFER OF VOLUME DATAS

Errors I encountered, what they mean, and how you fixed them

I list each error exactly as seen during your session (or very close), why it happened, and the exact command(s) you ran to fix it.

Error 1 — `tee: invalid option -- '/'`

Cause: you typed `tee -a/etc/fstab` without a space between `-a` and `/etc/fstab`.

Symptom / exact message:

`tee: invalid option -- '/'`

Fix: add a space between the `-a` flag and the path. Correct command:

```
echo 'UUID=<UUID> /data ext4 defaults,nofail 0 2' | sudo tee -a /etc/fstab
```

Why: `tee -a` takes an argument of a filename; without space it treated `/etc/fstab` as options.

Error 2 — `mount -a error: "mount point does not exist"`

Cause: `/etc/fstab` referenced `/data`, but `/data` folder did not exist at the time `mount -a` ran.

Symptom / message:

`mount: mount point /data does not exist`

Fix: create the directory, then re-run `mount -a`:

```
sudo mkdir -p /data
sudo mount -a
```

Note: after creating the directory you ran `mount -a` and it succeeded.

TASK 2 - AWS EBS _ SNAPSHOT _ TRANSFER OF VOLUME DATAS

Error 3 — systemd hint: “your fstab has been modified, but systemd still uses the old version”

Cause: /etc/fstab changed but systemd needed to reload its internal state.

Message:

hint: your fstab has been modified, but systemd still uses the old version
use 'systemctl daemon-reload' to reload

Fix: reload systemd then mount:

```
sudo systemctl daemon-reload
sudo mount -a
```

Alternate commands: sudo systemctl daemon-reexec (re-executes systemd process)
— both are fine.

Error 4 — df -h /data shows /dev/root (size ~6.8G) instead of 10G

Cause: the /data path was not actually mounted to the 10G EBS; instead /data was just an empty directory on the root filesystem.

Symptom: df -h /data displays /dev/root and a small size.

Fix / checks:

```
# Show block device info and what is mounted where
lsblk -f
mount | grep /data
# If the 10G device exists, mount it manually
sudo mount /dev/nvme1n1 /data
# Then confirm
df -h /data
```

Why: mount must succeed to attach the external 10G filesystem to /data. If mount failed earlier or fstab had a wrong UUID, /data will just be a directory on root disk.

Error 5 — “busy” when trying to unmount

Cause: some shell or process had a working directory or an open file on /data.

Fix: find processes and stop them, then unmount:

```
# find processes using /data
sudo lsof +f -- /data
# or
sudo fuser -m /data
# kill the offending PID(s) or close sessions, then
sudo umount /data Why: unmount fails while files are open on that filesystem.
```

TASK 2 - AWS EBS _ SNAPSHOT _ TRANSFER OF VOLUME DATAS

Debugging tips you used or can use

- `lsblk -f` — shows devices, filesystems, sizes and mount points.
 - `blkid` or `sudo blkid -s UUID -o value /dev/nvme1n1` — prints filesystem UUID.
 - `df -h /data` — shows what is mounted at /data and the apparent size.
 - `mount | grep /data` — confirms which device is mounted there.
 - `dmesg | tail -n 20` — kernel messages (good if mount fails or device not recognized).
 - `sudo systemctl daemon-reload` — reload systemd after fstab edits.
 - `sudo cp /etc/fstab /etc/fstab.bak.$(date +%F-%H%M)` — always take a backup of fstab before editing.
-

Command reference (short, one-line explanations)

- `lsblk -f` — list block devices and filesystems (useful to find the new EBS).
 - `sudo mkfs -t ext4 /dev/nvme1n1` — create ext4 filesystem (DESTROYS previous data on that device).
 - `sudo mkdir -p /data` — create mount point.
 - `sudo mount /dev/nvme1n1 /data` — mount the device at /data.
 - `sudo blkid -s UUID -o value /dev/nvme1n1` — get the UUID used for fstab.
 - `echo 'UUID=<UUID> /data ext4 defaults,nofail 0 2' | sudo tee -a /etc/fstab` — append fstab entry.
 - `sudo systemctl daemon-reload` — make systemd pick up fstab changes.
 - `sudo umount /data` — unmount before snapshot.
 - `sudo sync` — flush filesystem buffers to disk ahead of snapshot.
 - `aws ec2 create-snapshot --volume-id vol-xxxx` — (CLI) create snapshot; or do it from console.
 - `sudo lsof +f -- /data / sudo fuser -m /data` — find processes using the mount.
 - `sudo resize2fs /dev/nvme1n1` — expand ext4 filesystem if volume is grown.
-

Notes, gotchas, and best practices

- **AZ restriction:** Volumes can only be attached to instances in the same Availability Zone. When creating a volume from a snapshot, ensure you choose the AZ of the target instance.

TASK 2 - AWS EBS _ SNAPSHOT _ TRANSFER OF VOLUME DATAS


- **Unmount for consistency:** For the cleanest snapshot, unmount the filesystem before creating the snapshot. If you can't unmount, at least sync the filesystem and stop services writing to it.
 - **UUID vs device name:** Use UUID in `/etc/fstab` (UUID does not change across reboots or re-attachments), not `/dev/nvmeNn`.
 - **nofail in fstab:** avoids the instance failing to boot if a removable volume is missing.
 - **Backup / careful with mkfs:** `mkfs` will destroy data — double-check device name before running it.
 - **Charges:** snapshots and extra volumes cost money. Delete them when not needed.
-

Short summary of what I accomplished

1. Launched Instance-1 and attached a 10 GiB EBS volume.
2. Formatted and mounted that volume at `/data` on Instance-1 and made it persist via `/etc/fstab`.
3. Created test files and took a snapshot of the 10 GiB volume.
4. Detached the original volume from Instance-1 and created a new volume from the snapshot in Instance-2's AZ.
5. Launched Instance-2, attached the new volume, mounted it at `/data` and verified the files successfully moved.

TASK 2 - AWS EBS _ SNAPSHOT _ TRANSFER OF VOLUME DATAS

Pictorial Representation (ASCII Diagram)



— End of document —

TASK 2 - AWS EBS _ SNAPSHOT _ TRANSFER OF VOLUME DATAS

ATTACHED MY WORKS HERE ↓↓

```
systemctl restart unattended-upgrades.service

No containers need to be restarted.

User sessions running outdated binaries:
ubuntu @ session #2: sshd[1077,1189]
ubuntu @ session #4: sshd[1079,1289]
ubuntu @ user manager service: systemd[1083]

No VM guests are running outdated hypervisor (qemu) binaries on this host.
root@ip-172-31-23-59:~# systemctl status apache2
● apache2.service - The Apache HTTP Server
   Loaded: loaded (/usr/lib/systemd/system/apache2.service; enabled; preset: enabled)
   Active: active (running) since Sun 2025-08-17 11:44:51 UTC; 16s ago
     Docs: https://httpd.apache.org/docs/2.4/
   Main PID: 18209 (apache2)
     Tasks: 55 (limit: 1072)
   Memory: 5.3M (peak: 5.6M)
      CPU: 35ms
   CGroup: /system.slice/apache2.service
           └─18209 /usr/sbin/apache2 -k start
             └─18211 /usr/sbin/apache2 -k start
               └─18212 /usr/sbin/apache2 -k start

Aug 17 11:44:51 ip-172-31-23-59 systemd[1]: Starting apache2.service - The Apache HTTP Server...
Aug 17 11:44:51 ip-172-31-23-59 systemd[1]: Started apache2.service - The Apache HTTP Server.
root@ip-172-31-23-59:~# systemctl enable apache2
Synchronizing state of apache2.service with SysV service script with /usr/lib/systemd/systemd-sysv-install.
Executing: /usr/lib/systemd/systemd-sysv-install enable apache2
root@ip-172-31-23-59:~# cd /var/www/html
root@ip-172-31-23-59:/var/www/html# ls
index.html
root@ip-172-31-23-59:/var/www/html# vi index.html
root@ip-172-31-23-59:/var/www/html# cat index.html
hi mr. rishabh joshi sir good evening,

this is the process of creating the ebs service

-MOULI

root@ip-172-31-23-59:/var/www/html# S
```

The screenshot displays the AWS Management Console interface. On the left, the navigation menu includes Savings Plans, Reserved Instances, Dedicated Hosts, Capacity Reservations, Images, Elastic Block Store (selected), Network & Security, and Load Balancing. The main content area shows the 'Volumes (2)' page with a table of existing volumes:


Size	IOPS	Throughput	Snapshot ID	Source volume ID	Created	Availability Zone	Volume state
8 GiB	3000	125	snap-011b387...	-	2025/08/17 17:04 GMT+5:...	eu-north-1a	In-use
10 GiB	3000	125	-	-	2025/08/17 17:25 GMT+5:...	eu-north-1a	Available

Below the table, there is a 'Fault tolerance for all volumes in this Region' section and a 'Snapshot summary' showing '0 / 1' recently backed up volumes.

At the bottom, a terminal window is open, showing the output of the 'lsblk' command:

```
root@ip-172-31-23-59:~# lsblk
NAME        MAJ:MIN RM  SIZE RO TYPE MOUNTPOINTS
loop0        7:0    0   27.2M  1 loop /snap/amazon-ssm-agent/11320
loop1        7:1    0   73.9M  1 loop /snap/core22/1981
loop2        7:2    0   50.9M  1 loop /snap/snapd/24505
nvme0n1      259:0    0    8G   0 disk
└─nvme0n1p1  259:1    0    7G   0 part /
   └─nvme0n1p14 259:2    0    4M   0 part
      └─nvme0n1p15 259:3    0  106M  0 part /boot/efi
         └─nvme0n1p16 259:4    0  913M  0 part /boot
nvme1n1      259:5    0   10G   0 disk
```

TASK 2 - AWS EBS _ SNAPSHOT _ TRANSFER OF VOLUME DATAS



EC2 > Volumes > Create volume

Create volume [Info](#)

Create an Amazon EBS volume to attach to any EC2 instance in the same Availability Zone.

Volume settings

Volume type [Info](#)

General Purpose SSD (gp3)

Size (GiB) [Info](#)

10

Min: 1 GiB, Max: 16384 GiB.

IOPS [Info](#)

3000

Min: 3000 IOPS, Max: 16000 IOPS.

Throughput (MiB/s) [Info](#)

125


Min: 125 MiB, Max: 1000 MiB. Baseline: 125 MiB/s.

Availability Zone [Info](#)

eu-north-1a

Snapshot ID - optional [Info](#)

TASK 2 - AWS EBS _ SNAPSHOT _ TRANSFER OF VOLUME DATAS




A terminal window on an Ubuntu instance (IP: 51.20.66.119) displaying the output of the `lsblk` command. The output lists various block devices including loop devices and NVMe disks with their respective sizes, types, and mount points.

```
root@ip-172-31-23-59:~# lsblk
NAME        MAJ:MIN RM  SIZE RO  TYPE MOUNTPOINTS
loop0        7:0      0   27.2M  1 loop /snap/amazon-ssm-agent/11320
loop1        7:1      0   73.9M  1 loop /snap/core22/1981
loop2        7:2      0   50.9M  1 loop /snap/snapd/24505
nvme0n1     259:0    0    8G   0 disk 
  nvme0n1p1 259:1    0    7G   0 part /
  nvme0n1p14 259:2    0    4M   0 part 
  nvme0n1p15 259:3    0  106M  0 part /boot/efi
  nvme0n1p16 259:4    0   913M  0 part /boot
root@ip-172-31-23-59:~#
```

Snapshots (1) Info


Owned by me Last updated less than a minute ago [Recycle Bin](#) [Actions](#) [Create snapshot](#)

☐	Name	Snapshot ID	Full snapshot size	Volume size	Description	Storage tier	Snapshot status
☐		snap-0b02f5ed61d8d1b8a	236 MiB	10 GiB	linux 1 volume snapshot	Standard	Completed



A terminal window on an Ubuntu instance (IP: 51.20.66.119) showing the creation of a directory and a file, and then listing the contents of the directory.

```
root@ip-172-31-23-59:~# echo "hello from instance -1 " | tee /data/hello.txt
hello from instance -1
root@ip-172-31-23-59:~# mkdir -p /data/docs
root@ip-172-31-23-59:~# echo "this will move via snapshot" | tee /data/docs/readme.txt
this will move via snapshot
root@ip-172-31-23-59:~# ls -l /data/docs
total 4
-rw-r--r-- 1 root root 28 Aug 17 12:40 readme.txt
root@ip-172-31-23-59:~#
```



The AWS Management Console shows the 'Volumes' page with a table of volumes. A modal dialog is open, asking for confirmation to detach the volume `vol-0b93c41bb9e9f0a90`.

Name	Volume ID	Type	Size	IOPS	Throughput	Snapshot ID	Source volume ID	Created
	vol-06f606b9bfb2c5a3	gp3	8 GiB	3000	125	snap-011b387...	-	2025/...
☒	vol-0b93c41bb9e9f0a90	gp3	10 GiB	3000	125	-	-	2025/...

Detach vol-0b93c41bb9e9f0a90?

After you detach a volume, you might still be charged for volume storage. If you no longer need the volume, delete it to stop incurring charges.

Are you sure that you want to detach volume vol-0b93c41bb9e9f0a90?


[Cancel](#) [Detach](#)

Volume ID	Size	Type	Status check
vol-0b93c41bb9e9f0a90	10 GiB	gp3	Okay

AWS Compute Optimizer finding	Volume state	IOPS	Throughput
Opt-in to AWS Compute Optimizer for recommendations. Learn more	In-use	3000	125

Fast snapshot restored	Availability Zone	Created	Multi-Attach enabled
No	eu-north-1a	Sun Aug 17 2025 17:25:37 GMT+0530 (India Standard Time)	No

TASK 2 - AWS EBS _ SNAPSHOT _TRANSFER OF VOLUME DATAS



The screenshot shows a terminal window with a file manager on the left. The terminal output includes system statistics, update status, and disk details.

```
System load: 0.0      Temperature: -273.1 C
Usage of /: 25.4% of 6.71GB    Processes: 151
Memory usage: 3%      Users logged in: 0
Swap usage: 0%        IPv4 address for enp3s0: 172.31.27.107

Expanded Security Maintenance for Applications is not enabled.

0 updates can be applied immediately.

Enable ESM Apps to receive additional future security updates.
See https://ubuntu.com/esm or run: sudo pro status

The list of available updates is more than a week old.
To check for new updates run: sudo apt update

The programs included with the Ubuntu system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*/copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by
applicable law.

/usr/bin/xaauth: file /home/ubuntu/.Xauthority does not exist
To run a command as administrator (user "root"), use "sudo <command>".
See "man sudo_root" for details.

ubuntu@ip-172-31-27-107:~$ lsblk
NAME        MAJ:MIN RM  SIZE RO TYPE MOUNTPOINTS
loop0       7:0      0  27.2M  1 loop /snap/amazon-ssm-agent/11320
loop1       7:1      0  73.9M  1 loop /snap/core22/1981
loop2       7:2      0  50.9M  1 loop /snap/snapd/24505
nvme0n1     259:0    0    8G  0 disk 
├─nvme0n1p1 259:1    0    7G  0 part /
├─nvme0n1p14 259:2    0    4M  0 part 
├─nvme0n1p15 259:3    0  106M  0 part /boot/efi
├─nvme0n1p16 259:4    0  913M  0 part /boot
└─nvme1n1   259:5    0   10G  0 disk 
ubuntu@ip-172-31-27-107:~$
```


The programs included with the Ubuntu system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*/copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by
applicable law.

/usr/bin/xaauth: file /home/ubuntu/.Xauthority does not exist
To run a command as administrator (user "root"), use "sudo <command>".
See "man sudo_root" for details.

```
ubuntu@ip-172-31-27-107:~$ lsblk
NAME        MAJ:MIN RM  SIZE RO TYPE MOUNTPOINTS
loop0       7:0      0  27.2M  1 loop /snap/amazon-ssm-agent/11320
loop1       7:1      0  73.9M  1 loop /snap/core22/1981
loop2       7:2      0  50.9M  1 loop /snap/snapd/24505
nvme0n1     259:0    0    8G  0 disk 
├─nvme0n1p1 259:1    0    7G  0 part /
├─nvme0n1p14 259:2    0    4M  0 part 
├─nvme0n1p15 259:3    0  106M  0 part /boot/efi
├─nvme0n1p16 259:4    0  913M  0 part /boot
└─nvme1n1   259:5    0   10G  0 disk 
ubuntu@ip-172-31-27-107:~$ mkdir -p /data
mkdir: cannot create directory '/data': Permission denied
ubuntu@ip-172-31-27-107:~$ sudo -i
root@ip-172-31-27-107:~# mkdir -p /data
root@ip-172-31-27-107:~# mount /dev/nvme1n1 /data
root@ip-172-31-27-107:~# -h /data
-h: command not found
root@ip-172-31-27-107:~# df -h /data
Filesystem      Size  Used Avail Use% Mounted on
/dev/nvme1n1    9.8G  36K  9.3G   1% /data
root@ip-172-31-27-107:~# ls -l /data
total 24
drwxr-xr-x 2 root root 4096 Aug 17 12:39 docs
-rw-r--r-- 1 root root  24 Aug 17 12:39 hello.txt
drwx----- 2 root root 16384 Aug 17 12:01 lost+found
root@ip-172-31-27-107:~#
```

TASK 2 - AWS EBS _ SNAPSHOT _TRANSFER OF VOLUME DATAS



Instances (1/2) Info

Last updated less than a minute ago

Find Instance by attribute or tag (case-sensitive)

All states

	Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4
☐	linux1	i-0e50b51bf3096577c	Running	t3.micro	3/3 checks passed	View alarms +	eu-north-1a	ec2-51-20-6...
☒	linux 2	i-0ab772be786663618	Running	m7i-flex.large	3/3 checks passed	View alarms +	eu-north-1a	ec2-13-60-2...

Snapshots (1) Info

Last updated less than a minute ago

Owned by me

Search

Recycle Bin

Actions

Create snapshot

	Name	Snapshot ID	Full snapshot size	Volume size	Description	Storage tier	Snapshot status
☐		snap-0b02f5ed61d8d1b8a	236 MiB	10 GiB	linux 1 volume snapshot	Standard	Completed

Volumes (3) Info

Last updated less than a minute ago

Actions

Create volume

Saved filter sets

Choose filter set

Search

	Name	Volume ID	Type	Size	IOPS	Throughput	Snapshot ID	Source volume ID	Created
☐	changing vol	vol-0c5b2ed68317e7d13	gp3	10 GiB	3000	125	snap-0b02f5e...	-	2025/
☐	linux 1 vol	vol-06f606b9bf8b2c5a3	gp3	8 GiB	3000	125	snap-011b387...	-	2025/
☐	linux 2 vol	vol-0f588e103b57db92f	gp3	8 GiB	3000	125	snap-011b387...	-	2025/